

Computer Vision Assignment 01: Introduction to PyTorch Report

Computer Vision Assignment 01: Introduction to PyTorch Report

2 Simple 2D Classifier

2.3 Training loop

2.4 Multi-layer perceptron

2.5 Feature transform

3 Digit classifier

3.3 Multi-layer perceptron

3.4 Convolutional network

3.5 Comparison of number of parameters

3.6 Confusion matrix

2 Simple 2D Classifier

2.3 Training loop

Question: Once this is done, you should be able to run the `train.py` script. What accuracy do you achieve with the linear classifier? Is this an expected result? Justify your answer.

```
>>> python train.py
2 [Epoch 01] Acc.: 43.0556%
3 [Epoch 02] Loss: 0.7339
  [Epoch 02] Acc.: 40.6746%
5 [Epoch 03] Loss: 0.7000
6 [Epoch 03] Acc.: 49.8016%
7 [Epoch 04] Loss: 0.6960
8 [Epoch 04] Acc.: 32.9365%
9 [Epoch 05] Loss: 0.6942
10 [Epoch 05] Acc.: 40.6746%
11 [Epoch 06] Loss: 0.6934
12 [Epoch 06] Acc.: 41.8651%
13 [Epoch 07] Loss: 0.6933
14 [Epoch 07] Acc.: 44.6429%
15 [Epoch 08] Loss: 0.6933
16 [Epoch 08] Acc.: 47.0238%
17 [Epoch 09] Loss: 0.6932
18 [Epoch 09] Acc.: 46.6270%
19 [Epoch 10] Loss: 0.6933
20 [Epoch 10] Acc.: 48.6111%
```

- The accuracy is `48.6111%` which is lower than expected (`50%` for random guess). However, the data shown below is not linearly separable, so the accuracy closed to `50%` becomes **reasonable**.

2.4 Multi-layer perceptron

Question: Switch to the new network by uncommenting `L83` in `train.py`. What accuracy does this network obtain? Why are the results better compared to the previous classifier?

```
1 >>> python train.py
2 [Epoch 01] Loss: 0.6409
3 [Epoch 01] Acc.: 51.3889%
4 [Epoch 02] Loss: 0.5039
5 [Epoch 02] Acc.: 95.2381%
6 [Epoch 03] Loss: 0.2507
7 [Epoch 03] Acc.: 99.4048%
8 [Epoch 04] Loss: 0.0961
9 [Epoch 04] Acc.: 99.6032%
10 [Epoch 05] Loss: 0.0484
11 [Epoch 05] Acc.: 99.6032%
12 [Epoch 06] Loss: 0.0299
13 [Epoch 06] Acc.: 99.6032%
14 [Epoch 07] Loss: 0.0210
15 [Epoch 07] Acc.: 99.6032%
16 [Epoch 08] Loss: 0.0160
17 [Epoch 08] Acc.: 99.6032%
18 [Epoch 09] Loss: 0.0128
19 [Epoch 09] Acc.: 99.8016%
20 [Epoch 10] Loss: 0.0111
21 [Epoch 10] Acc.: 99.8016%
```

- The final accuracy is `99.8016%`, which is better than the linear classifier.

2.5 Feature transform

Question: Think of a coordinate system that renders the two classes linearly separable and justify your choice.




Figure 1: Training data for 2D classifier.

- By observation, the data is not linearly separable since it is a circle. However, by changing the coordinate system to polar system or simply use $(x, y) \mapsto (x^2, y^2)$. Then the data will distribute on two lines:

$$\begin{aligned}x^2 + y^2 &= r_1 \\x^2 + y^2 &= r_2\end{aligned}\tag{1}$$

Therefore, the classes become linearly separable.

Question: Verify the hypothesis by training a linear classifier on the new representation.

```
>>> python train.py
```



```
15 [Epoch 05] Loss: 0.3092
16 [Epoch 05] Acc.: 91.2600%
```

Question: Next, use an MLP with one hidden layer of dimension 32 followed by ReLU and then the final linear prediction layer. What is the new testing accuracy?

```
1 self.layers = nn.Sequential(
2     nn.Linear(784, 32),
3     nn.ReLU(),
4     nn.Linear(32, 10)
5 )
```

```
1 >>> python train.py
2 100%|██████████| 3750/3750 [00:32<00:00, 114.94it/s]
3 [Epoch 01] Loss: 0.4166
4 [Epoch 01] Acc.: 91.4300%
5 100%|██████████| 3750/3750 [00:32<00:00, 117.09it/s]
6 [Epoch 02] Loss: 0.2859
7 [Epoch 02] Acc.: 92.6200%
8 100%|██████████| 3750/3750 [00:33<00:00, 113.34it/s]
9 [Epoch 03] Loss: 0.2429
10 [Epoch 03] Acc.: 93.5000%
11 100%|██████████| 3750/3750 [00:31<00:00, 117.46it/s]
12 [Epoch 04] Loss: 0.2165
13 [Epoch 04] Acc.: 94.3700%
14 100%|██████████| 3750/3750 [00:32<00:00, 115.52it/s]
15 [Epoch 05] Loss: 0.1986
16 [Epoch 05] Acc.: 93.6400%
```

3.4 Convolutional network

Question: What testing accuracy do you obtain with this architecture?

```
1 100%|██████████| 3750/3750 [00:42<00:00, 88.20it/s]
2 [Epoch 01] Loss: 0.2677
3 [Epoch 01] Acc.: 97.3400%
4 100%|██████████| 3750/3750 [00:41<00:00, 89.81it/s]
5 [Epoch 02] Loss: 0.0810
6 [Epoch 02] Acc.: 97.7400%
7 100%|██████████| 3750/3750 [00:49<00:00, 76.11it/s]
8 [Epoch 03] Loss: 0.0614
9 [Epoch 03] Acc.: 98.4200%
10 100%|██████████| 3750/3750 [00:41<00:00, 89.31it/s]
11 [Epoch 04] Loss: 0.0495
12 [Epoch 04] Acc.: 98.3700%
13 100%|██████████| 3750/3750 [00:41<00:00, 89.70it/s]
14 [Epoch 05] Loss: 0.0440
15 [Epoch 05] Acc.: 98.4500%
```

3.5 Comparison of number of parameters

Question: Compute the number of parameters of the MLP with one hidden layer. Compute the number of parameters of the convolutional network. **You should count both the weights and biases.** Provide detailed explanations and computations.

- MLP: For each layer, number of weight is the number of input dimension times the number of output dimension, number of bias is the output dimension, therefore:

$$Parameter_Num = \sum_i^{layers} (in_channel_i + 1)(out_channel_i) \quad (2)$$

- Convolution: For each layer, number of weight equals to the number of kernel times the number of weight in each kernel. Number of bias is the output channel number:


$$Parameter_Num = \sum_i^{layers} (in_channel_i \times width_i \times height_i + 1)(out_channel_i) \quad (3)$$

- The result is shown as below:

$$\begin{aligned} Parameter_Num_{MLP} &= (784 + 1) \times 32 + (32 + 1) \times 10 = 25450 \\ Parameter_Num_{convolution} &= 8 \times (1 + 3 \times 3 \times 1) + 16 \times (1 + 3 \times 3 \times 8) \\ &\quad + 32 \times (1 + 3 \times 3 \times 16) = 5888 \end{aligned} \quad (4)$$

3.6 Confusion matrix

Question: The confusion matrix $M_{i,j}$ is a useful tool for understanding the performance of a classifier. It is defined as follows: $M_{i,j}$ is the number of test samples for which the network predicted i , but the ground-truth label was j . Ideally, in the case of 100% accuracy, this will be a diagonal matrix. Based on the run validation epoch function from `train.py`, update `plot_confusion_matrix.py` to compute the confusion matrix. Provide the plot in the report and comment the results.



- The summation of the matrix equals to the size of the dataset. The performance of the model is good because the matrix is very closed to a diagonal matrix.